

# BIG DATA


Sirojul Munir | [rojulman@nurulfikri.ac.id](mailto:rojulman@nurulfikri.ac.id) | @rojulman

# Map Reduce

Sirojul Munir | [rojulman@nurulfikri.ac.id](mailto:rojulman@nurulfikri.ac.id) | @rojulman

# Latar Belakang

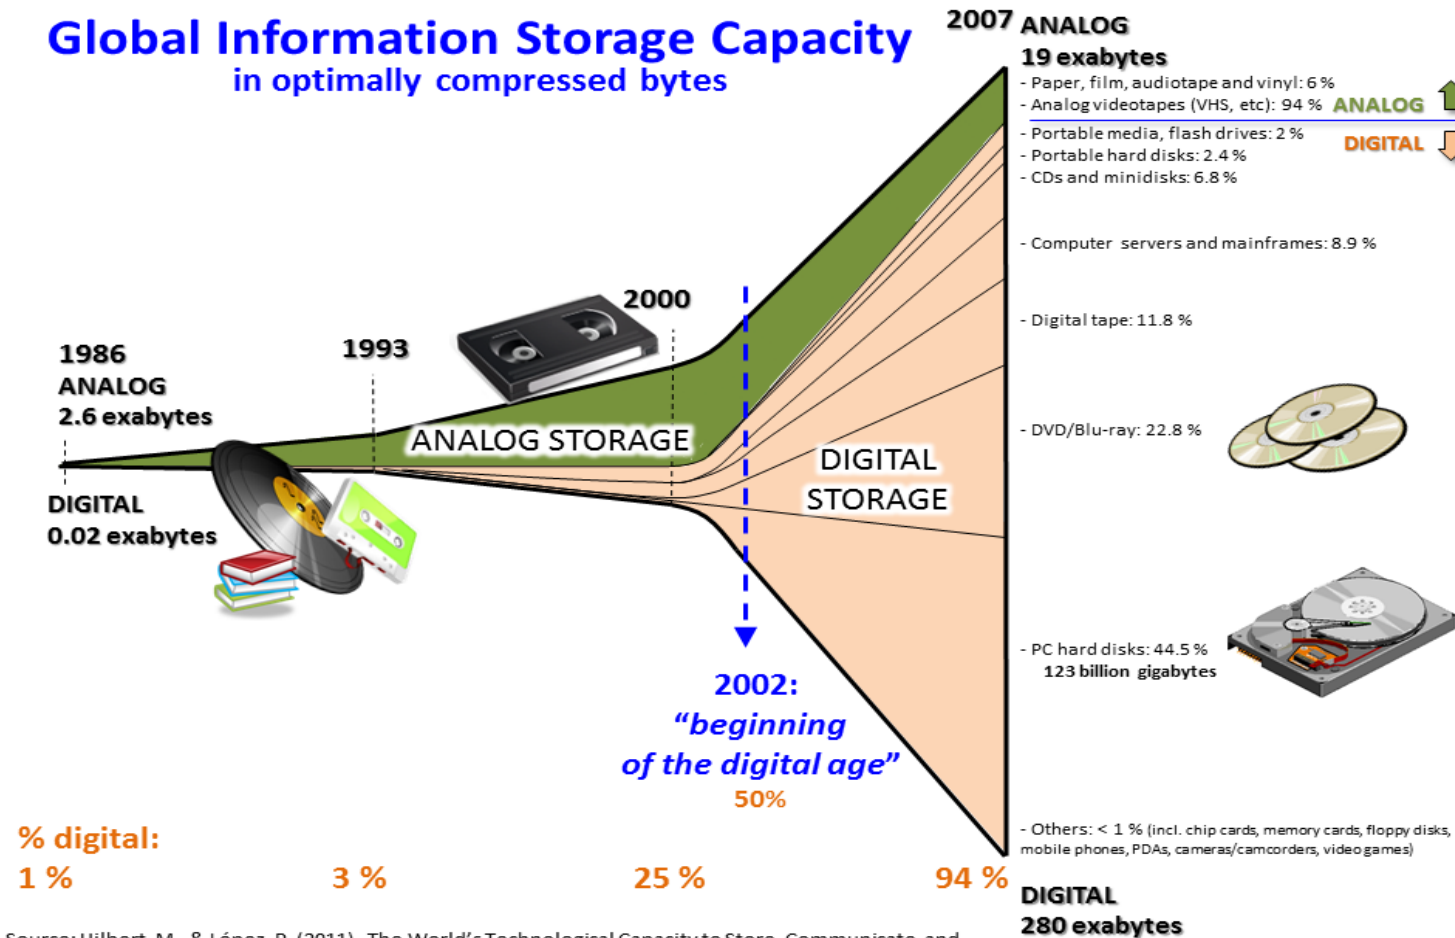
---



- Aplikasi tradisional melakukan **proses datanya pada sistem terpusat**
- Sistem terpusat tidak cocok untuk memproses data ber-skala besar karena **tidak dapat di akomodir oleh sistem standard database (RDBMS)**
- Sistem terpusat memiliki kelemahan (**adanya bottleneck**) ketika memproses multiple file secara simultan

# Latar Belakang : Growing Data

## Global Information Storage Capacity in optimally compressed bytes




Source: Hilbert, M., & López, P. (2011). The World's Technological Capacity to Store, Communicate, and Compute Information. *Science*, 332(6025), 60–65. <http://www.martinhilbert.net/WorldInfoCapacity.html>

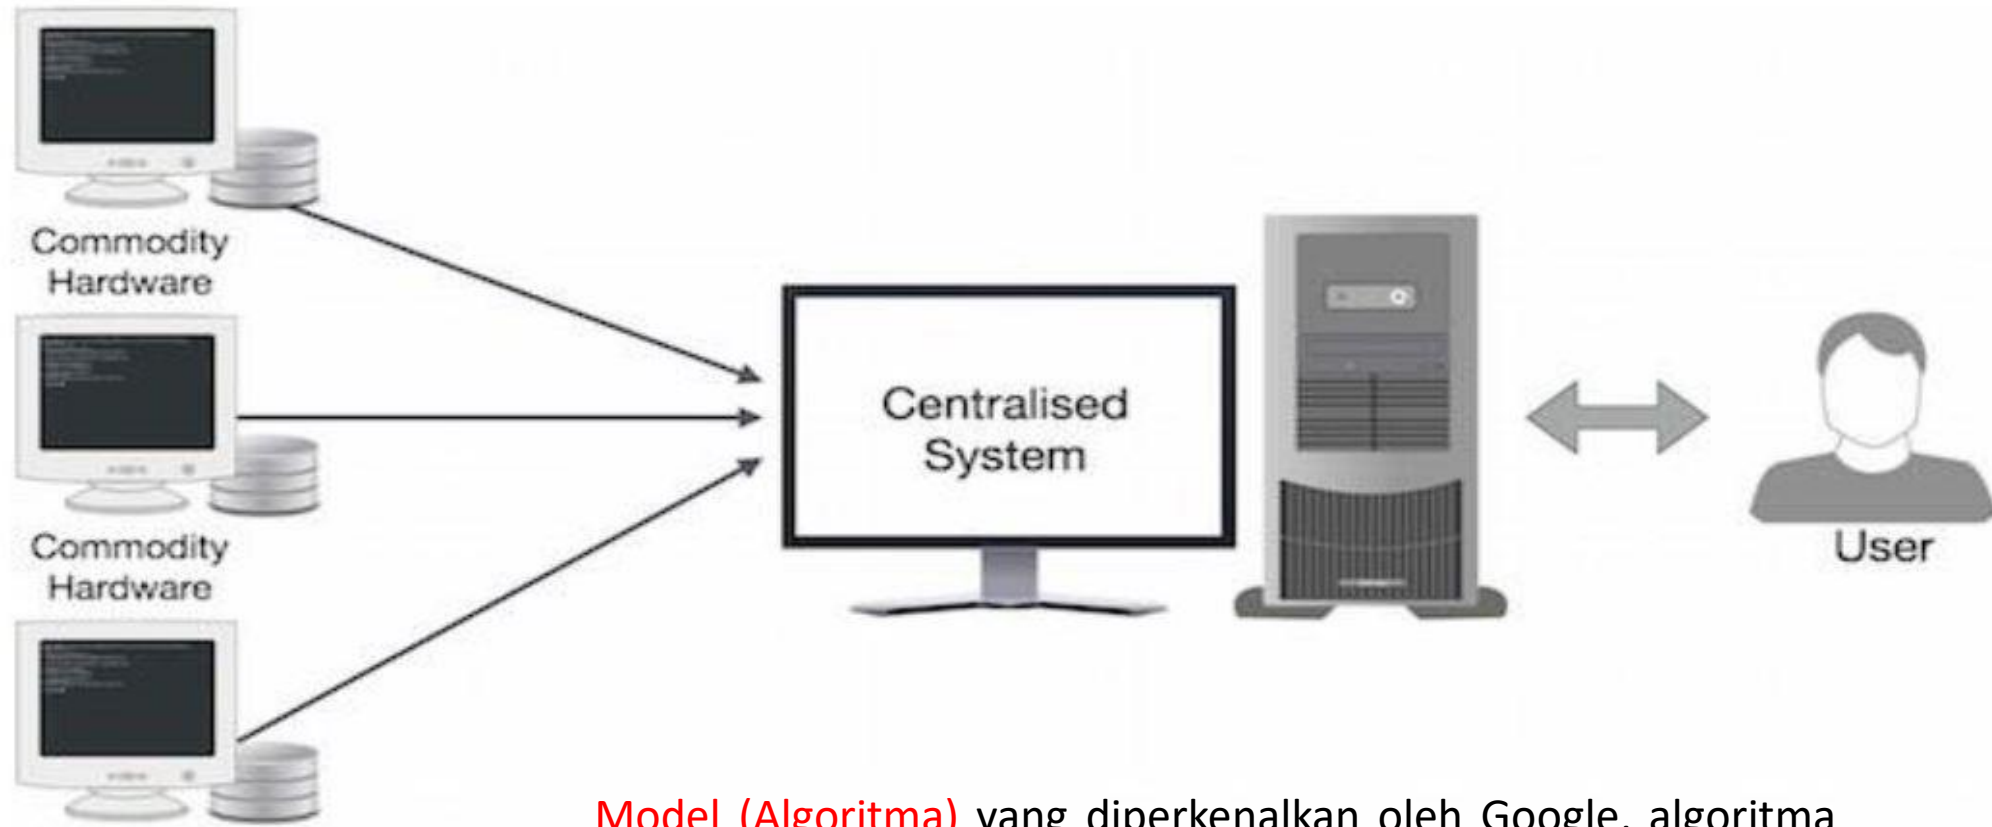
- Tahun 2009, Google memproses data per hari mencapai **24 Petabytes**
- Satu mesin server tidak dapat menangani seluruh proses data, karenanya dibutuhkan **system yang terdistribusi** dan **proses data yang paralel**

# Apa itu Map Reduce ?

- **Model (Algoritma)** yang diperkenalkan oleh Google, algoritma membagi pekerjaan (task) menjadi beberapa bagian pekerjaan dan dikerjakan oleh beberapa komputer
- Algoritma Map Reduce memecahkan masalah bottleneck yang terjadi pada pemrosesan data
- Proses data berskala besar:
  - Membutuhkan perangkat computer spek besar/tinggi
  - Eksekusi proses data secara terdistribusi
  - Menawarkan tingkat ketersediaan yang tinggi (high availability)

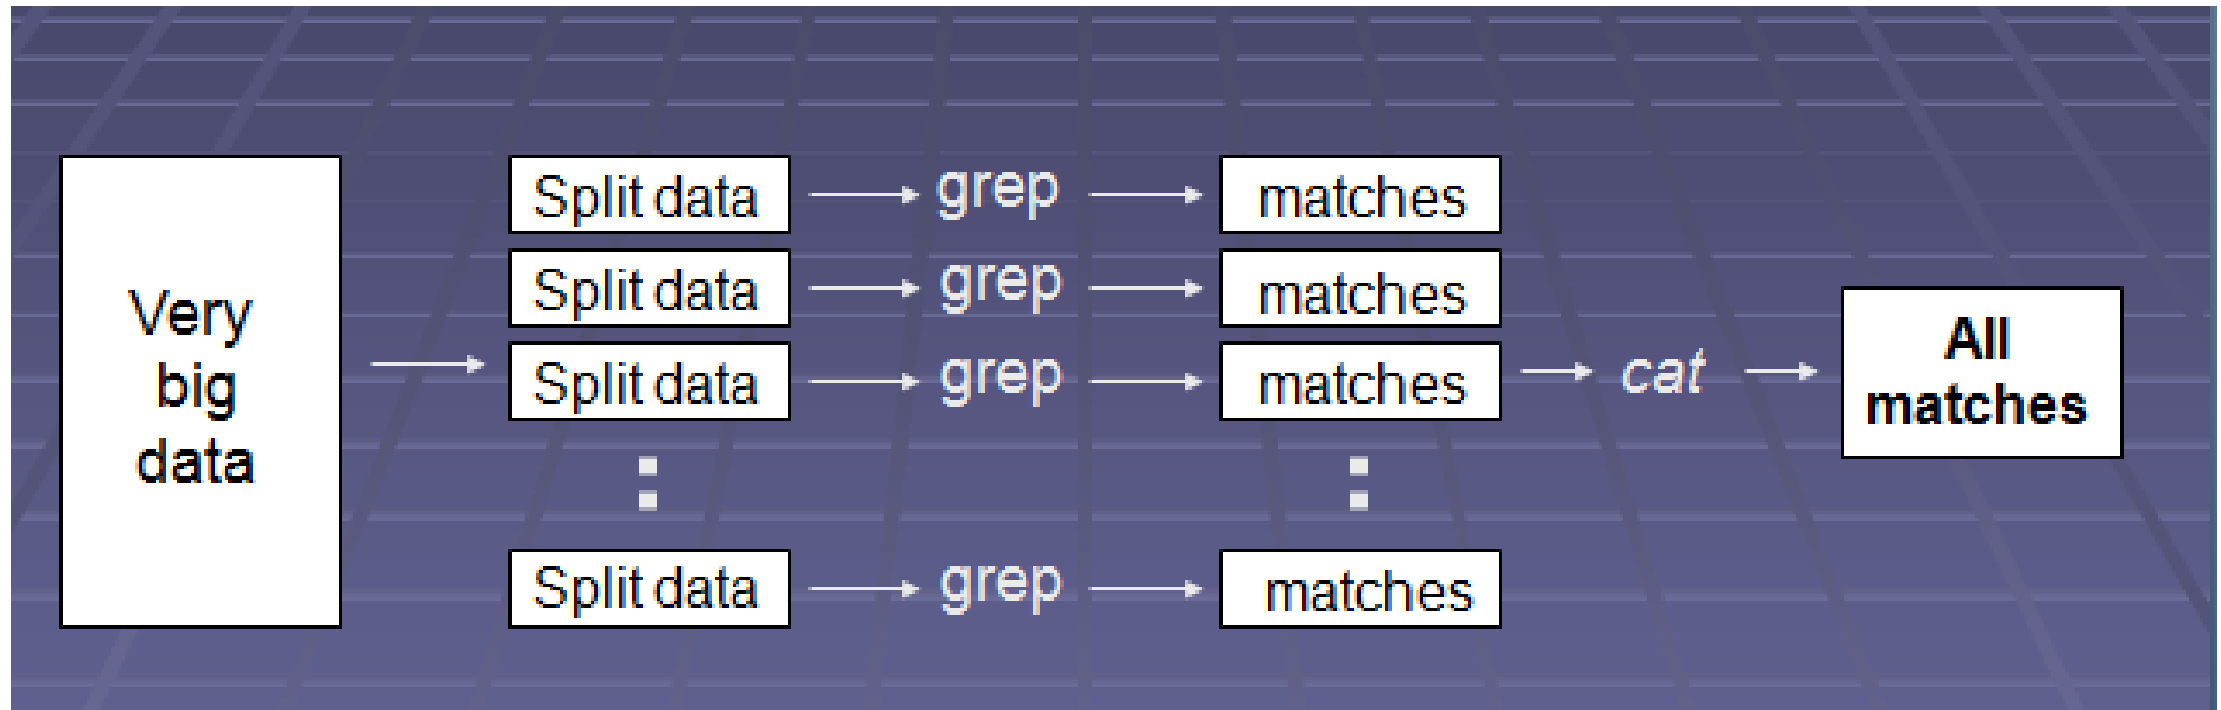


# Algoritma Mapreduce



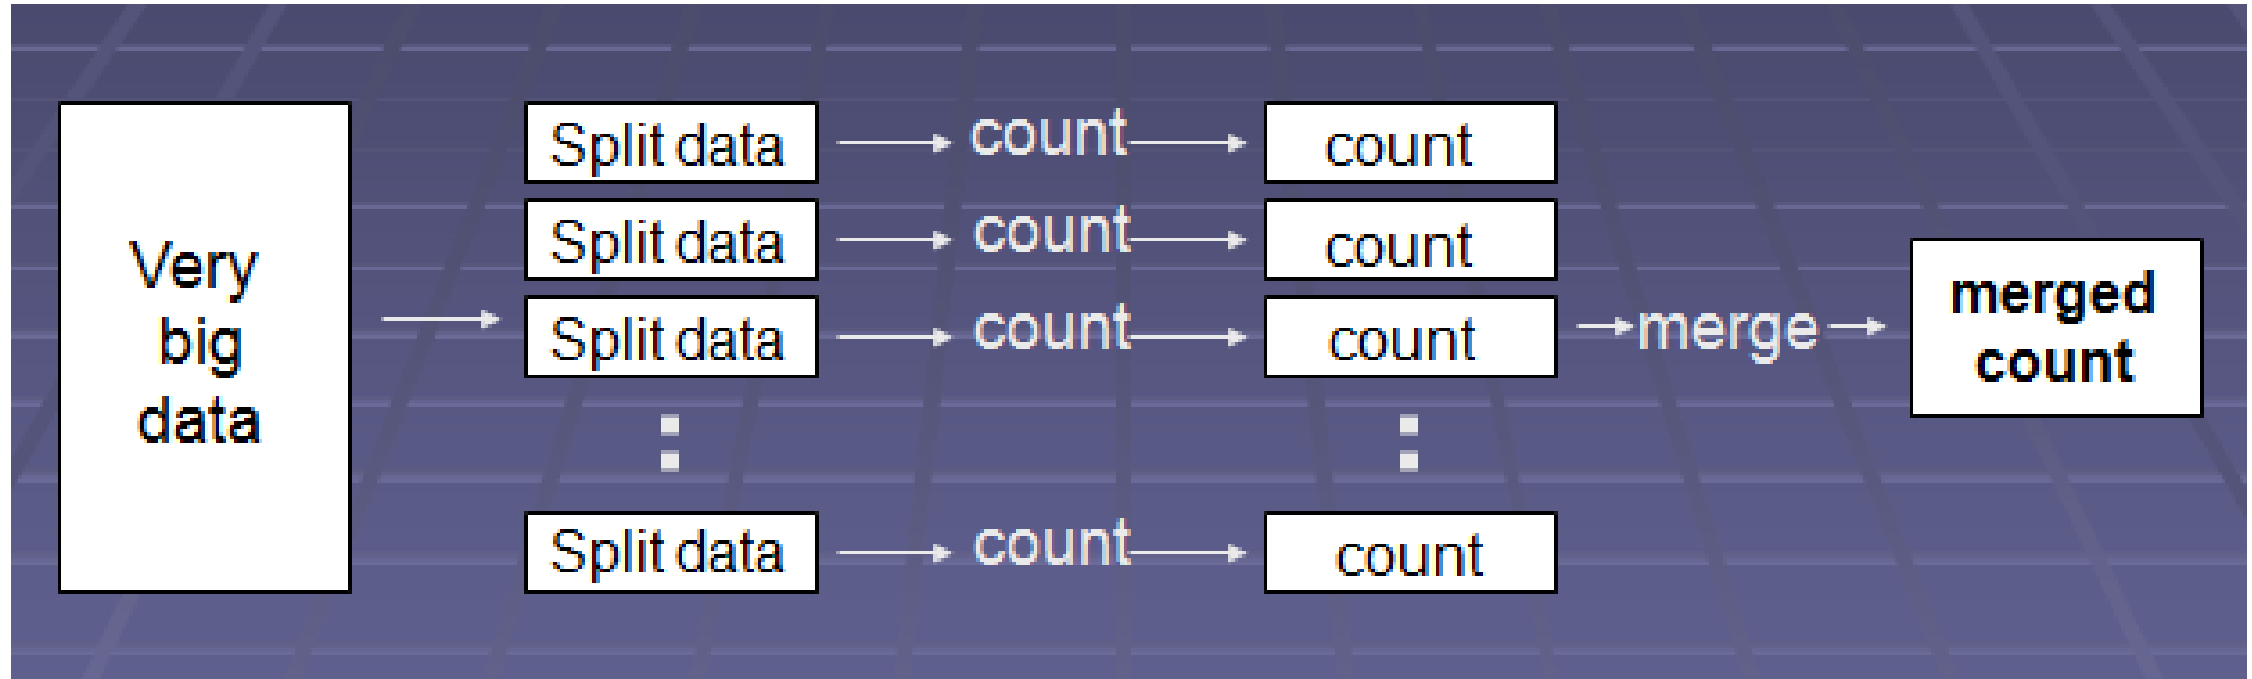
**Model (Algoritma)** yang diperkenalkan oleh Google, algoritma membagi pekerjaan (task) menjadi beberapa bagian pekerjaan dan dikerjakan oleh beberapa komputer

# Contoh : Proses data - Distributed Grep



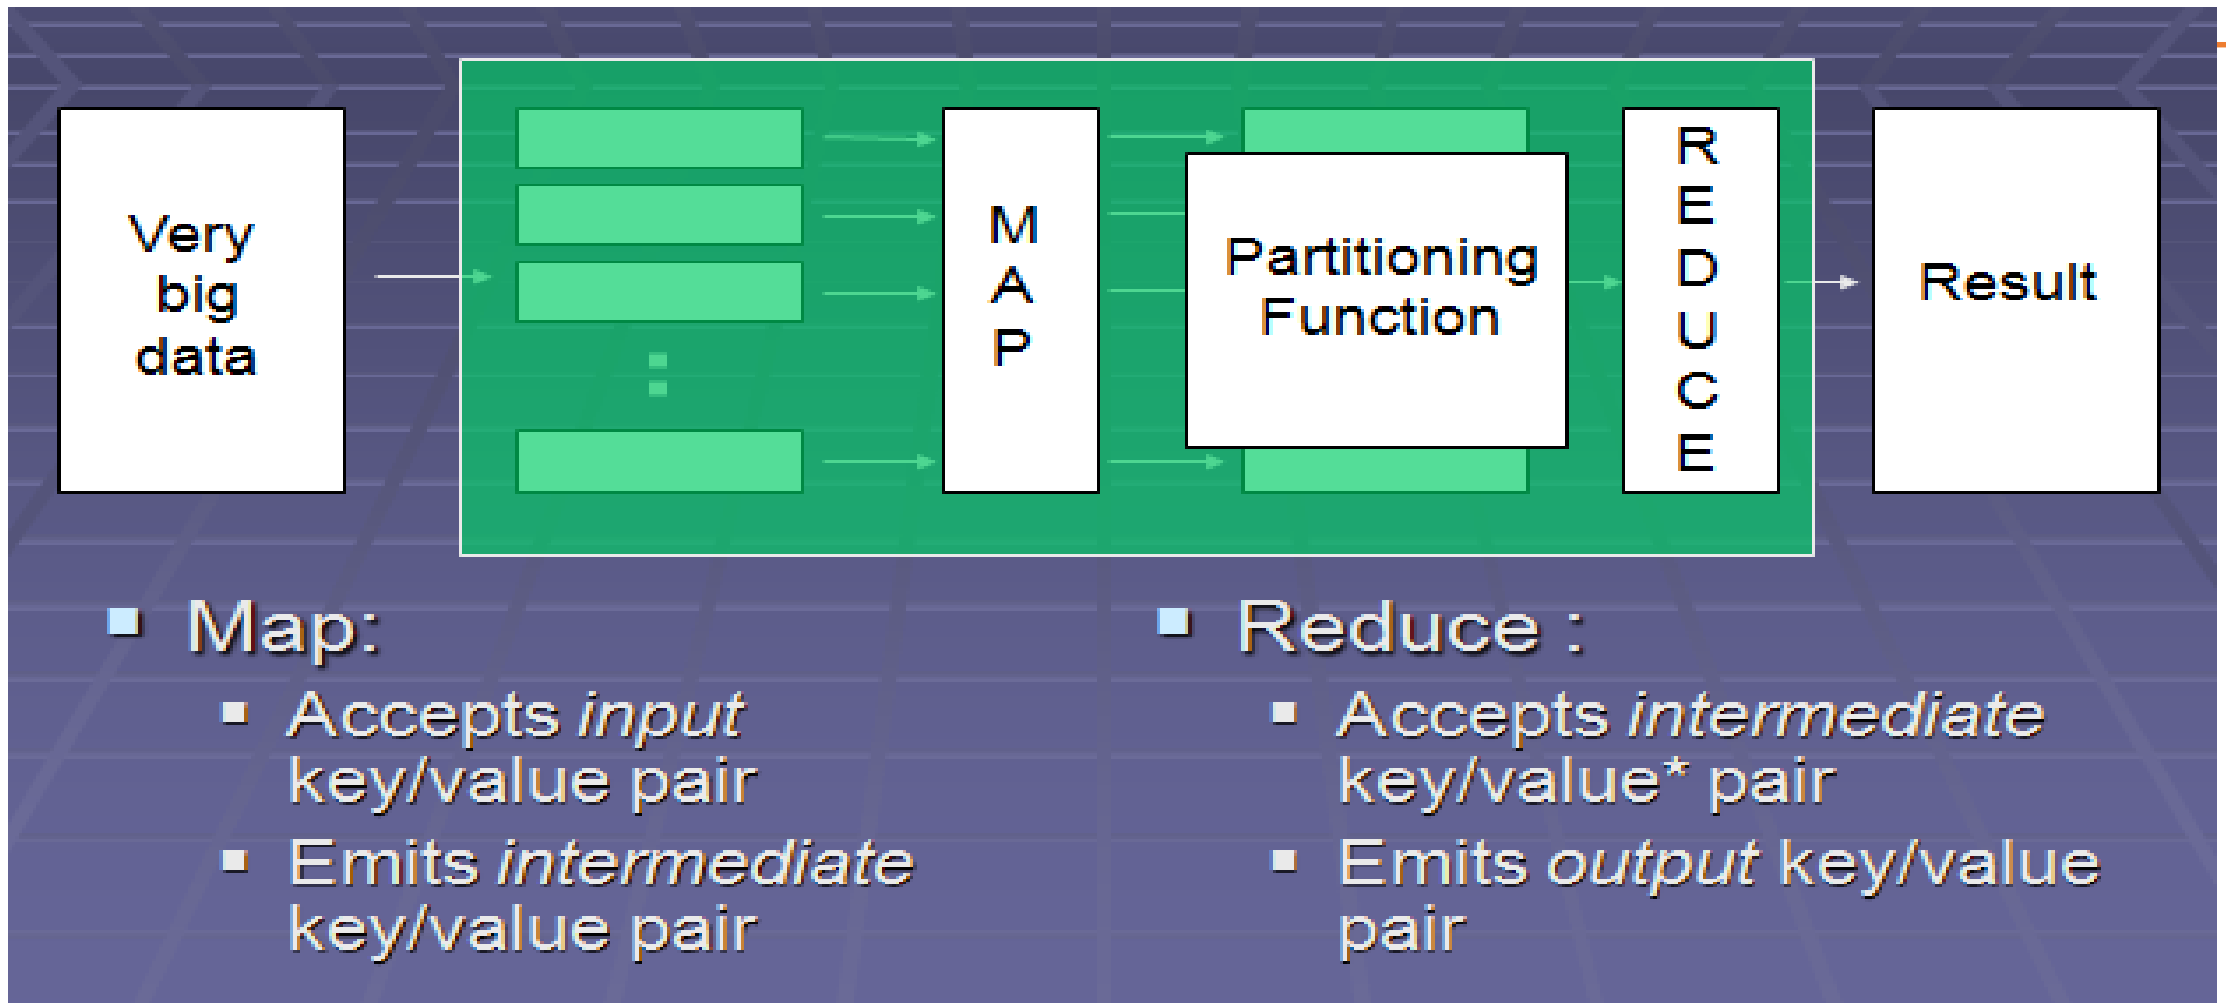
# Contoh : Proses data - Distributed Word Count

---


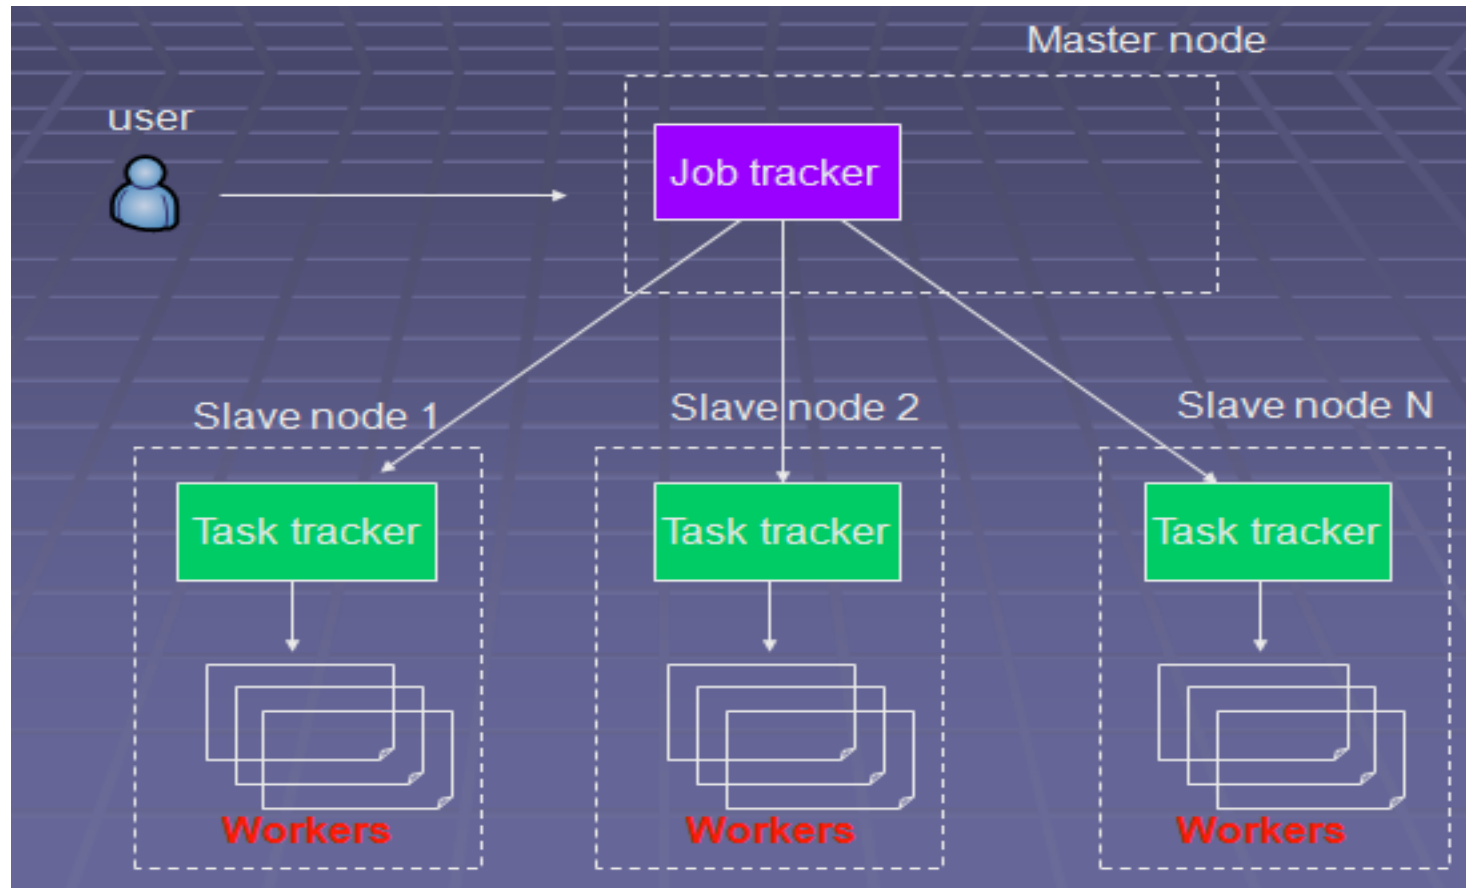




# Map + Reduce

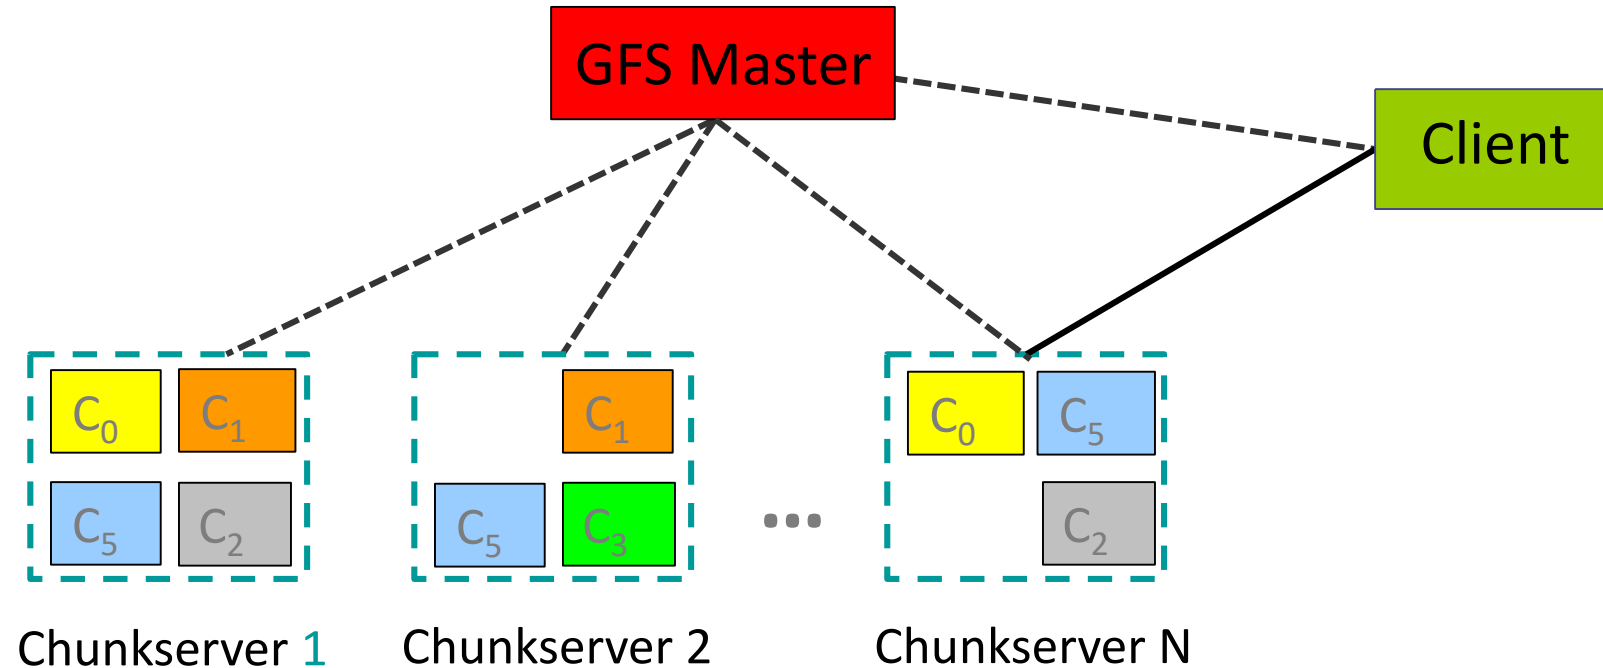


# Design System : Arsitektur view



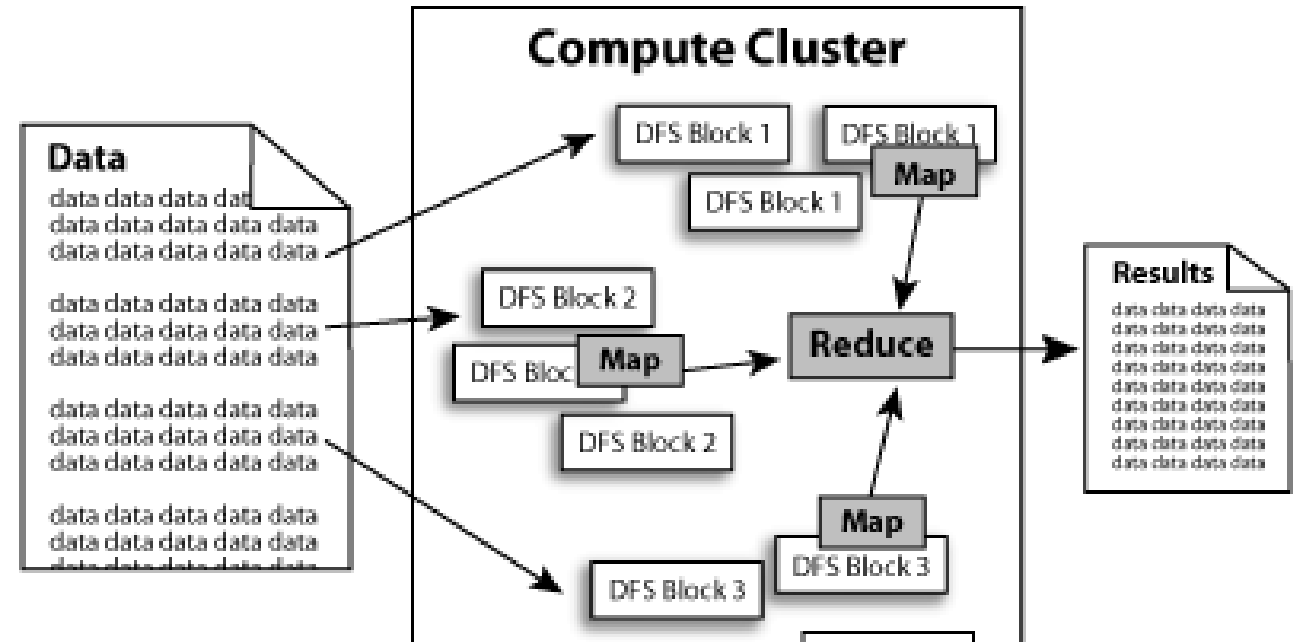
# Google File System (GFS)

- Goal
  - global view
  - make huge files available in the face of node failures
- Master Node (meta server)
  - Centralized, index all chunks on data servers
- Chunk server (data server)
  - File is split into contiguous chunks, typically 16-64MB.
  - Each chunk replicated (usually 2x or 3x).
    - Try to keep replicas in different racks.



# Model : Map Reduce

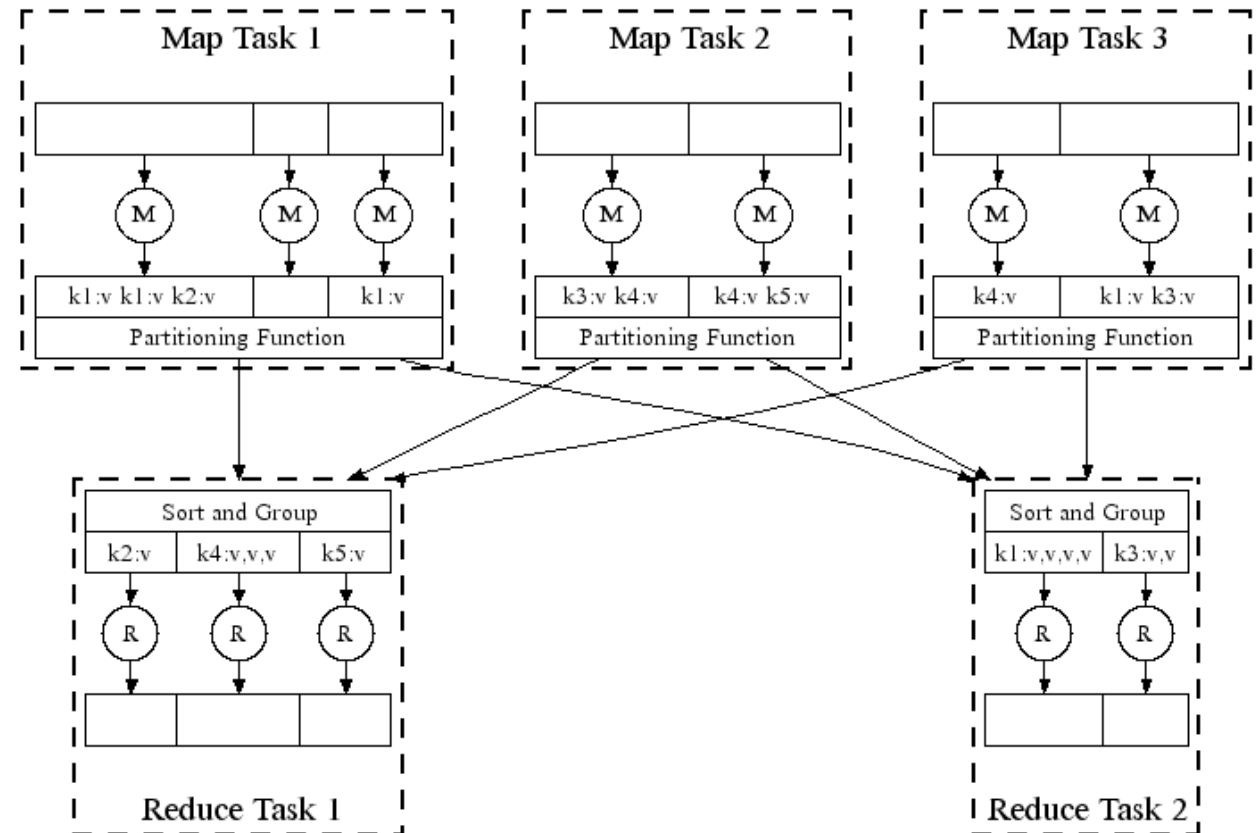
- Map
  - Process a key/value pair to generate intermediate key/value pairs
- Reduce
  - Merge all intermediate values associated with the same key
- Partition
  - By default :  $\text{hash}(\text{key}) \bmod R$
  - Well balanced



# Parallel Execution

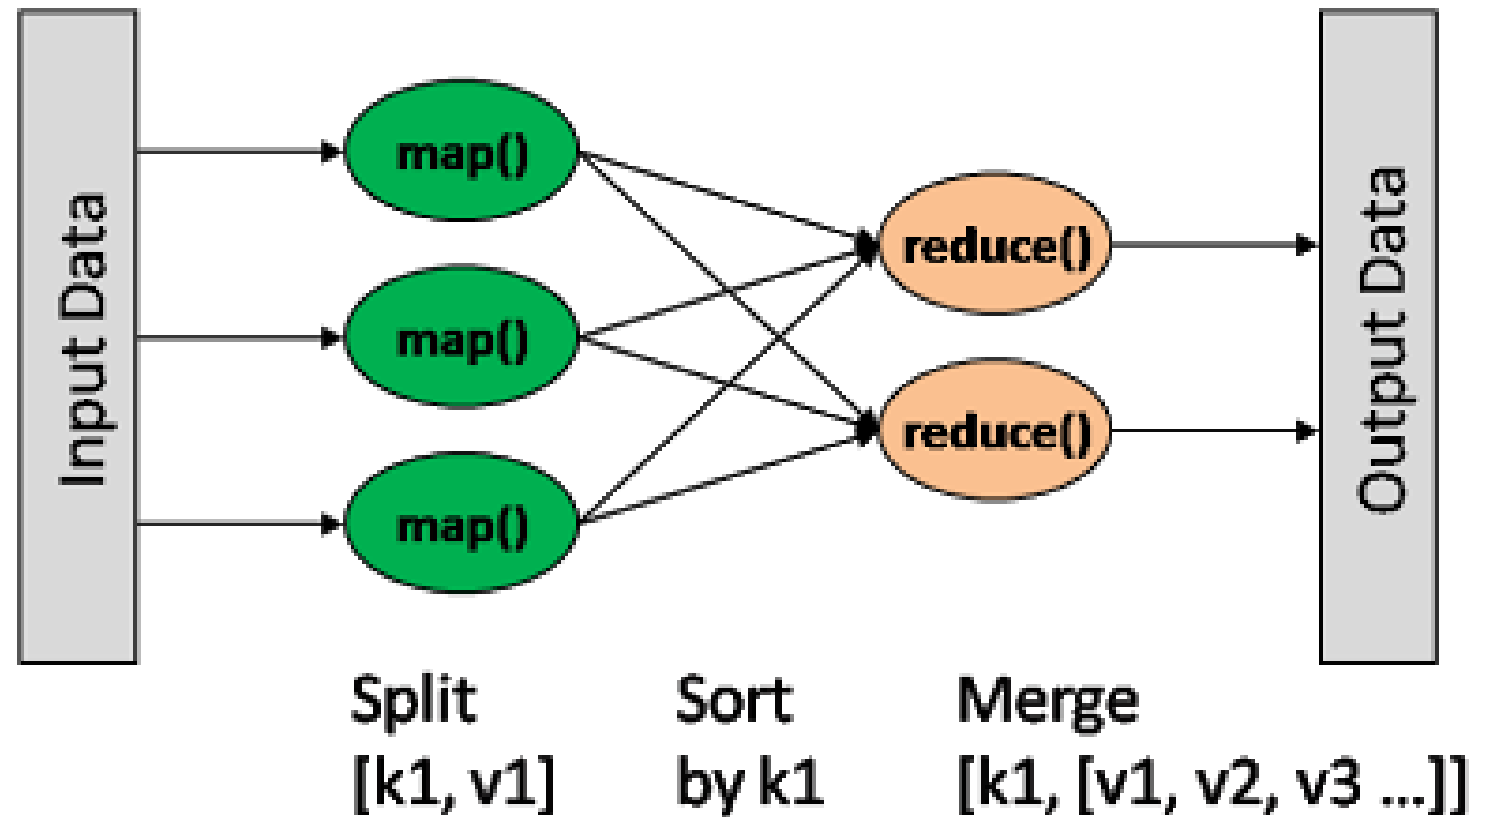
## Parallel Execution

- Map
  - Process a key/value pair to generate intermediate key/value pairs
- Reduce
  - Merge all intermediate values associated with the same key
- Partition
  - By default :  $\text{hash}(\text{key}) \bmod R$
  - Well balanced



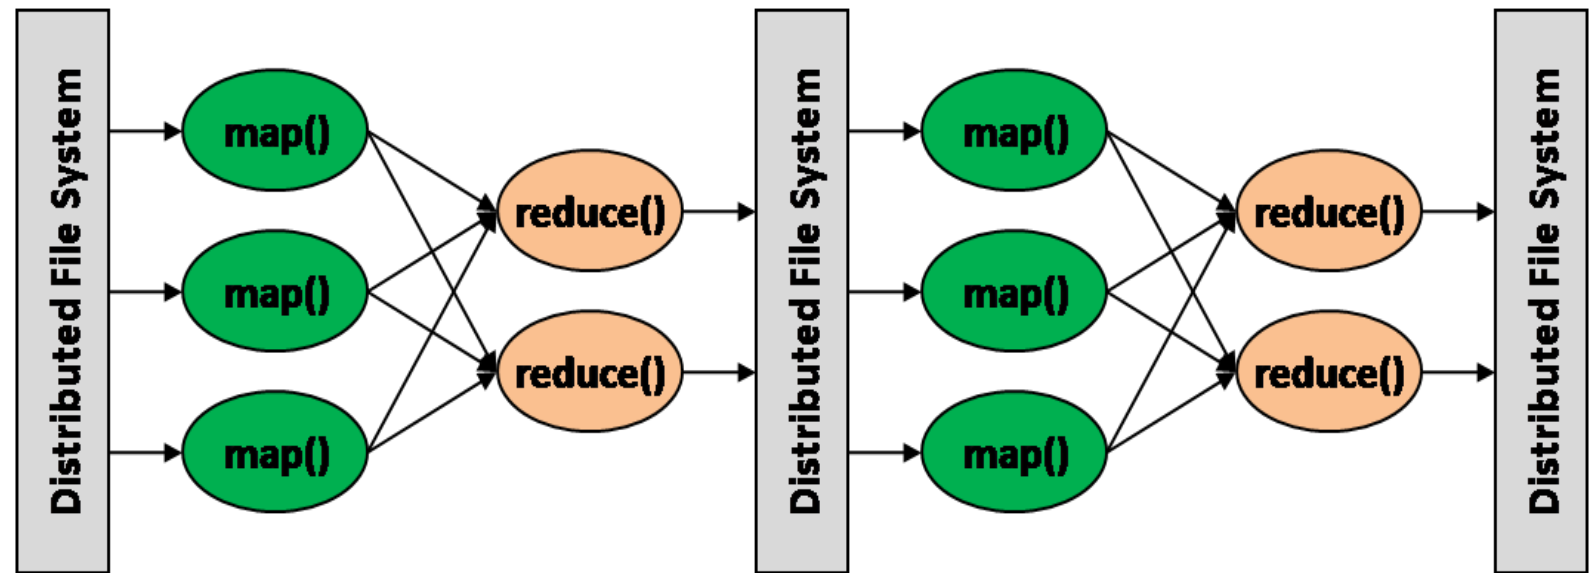
# Bagaimana Mapreduce bekerja ?

- **Map** – Membagi himpunan data menjadi beberapa himpunan data yang elemennya terdiri atas elemen key berpasangan dengan elemen value
- **Reduce** – Mengambil output dari map sebagai inputan dan melakukan kombinasi dari data tuple(key-value) menjadi bagian-bagian kecil data tuple




# Bagaimana Mapreduce bekerja ?

- **Map** – Membagi himpunan data menjadi beberapa himpunan data yang elemennya terdiri atas elemen key berpasangan dengan elemen value
- **Reduce** – Mengambil output dari map sebagai inputan dan melakukan kombinasi dari data tuple(key-value) menjadi bagian-bagian kecil data tuple



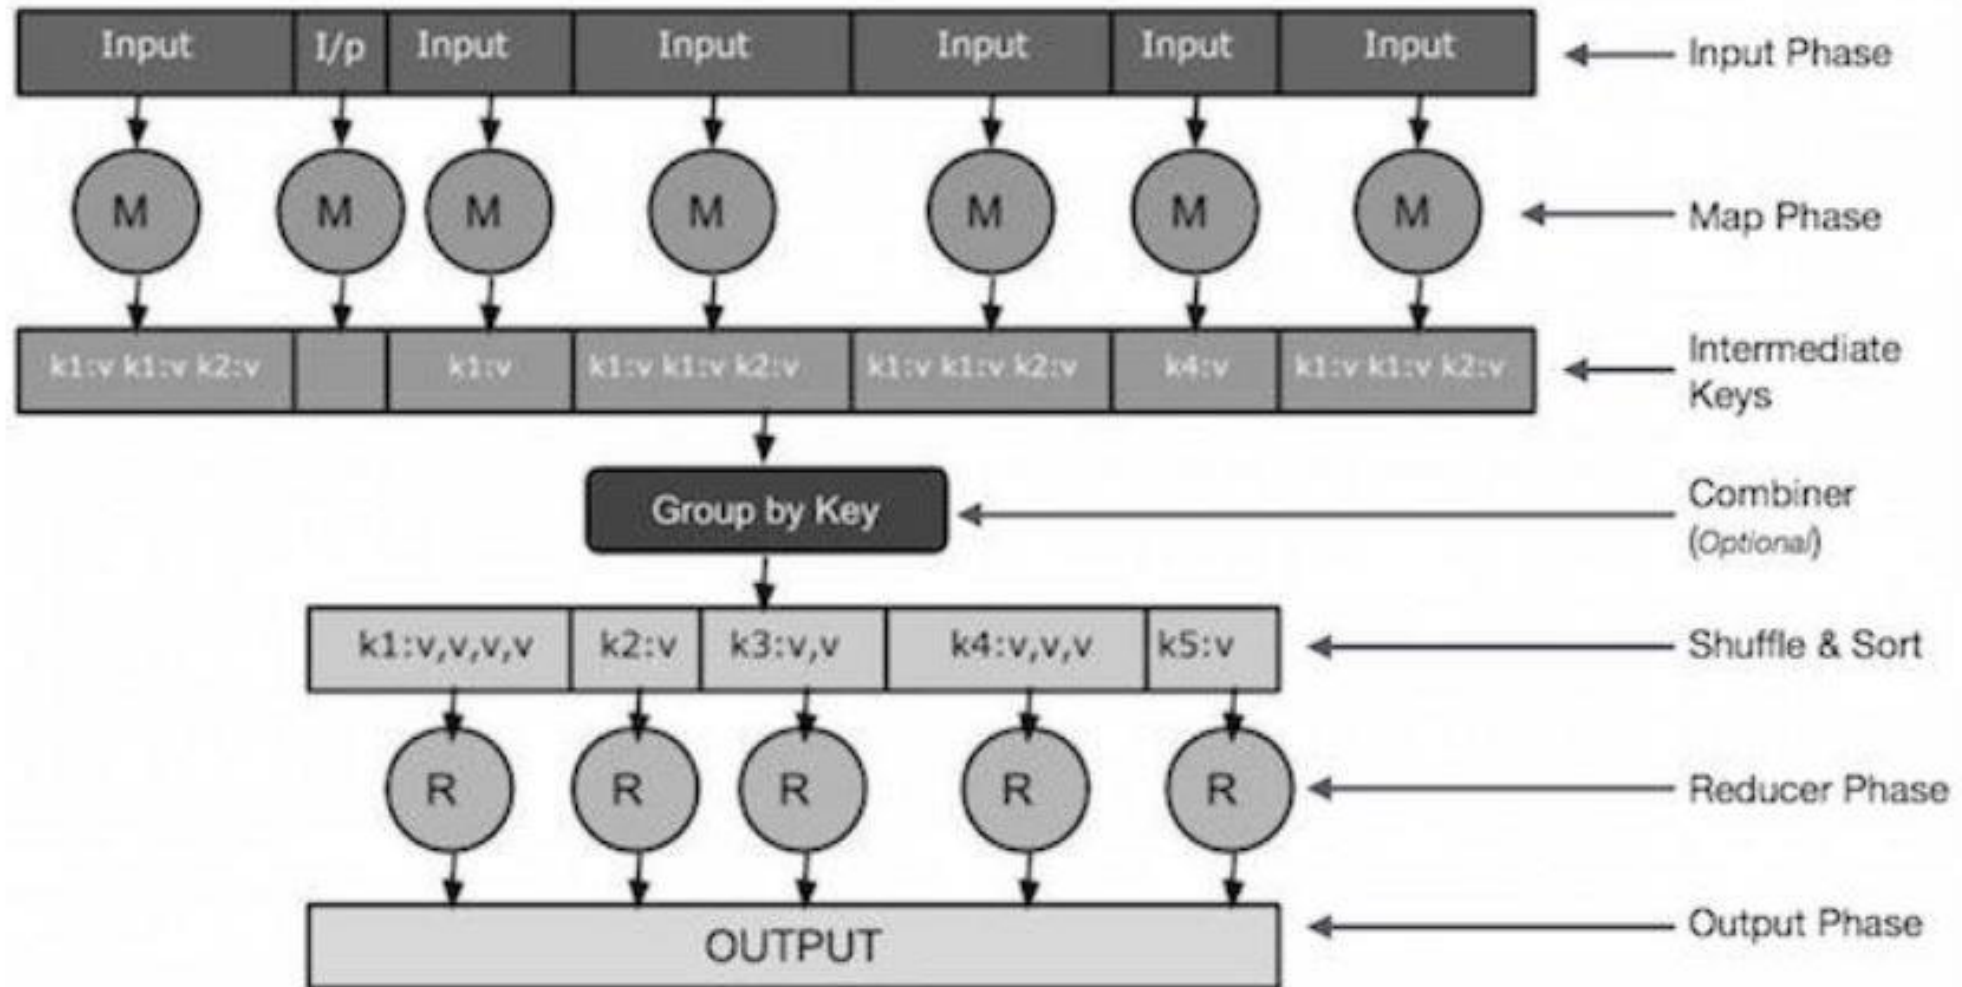
# Map - Reduce

	Input	Output
<b>Map</b>	$\langle k1, v1 \rangle$	list ( $\langle k2, v2 \rangle$ )
<b>Reduce</b>	$\langle k2, \text{list}(v2) \rangle$	list ( $\langle k3, v3 \rangle$ )





# Fase proses MapReduce

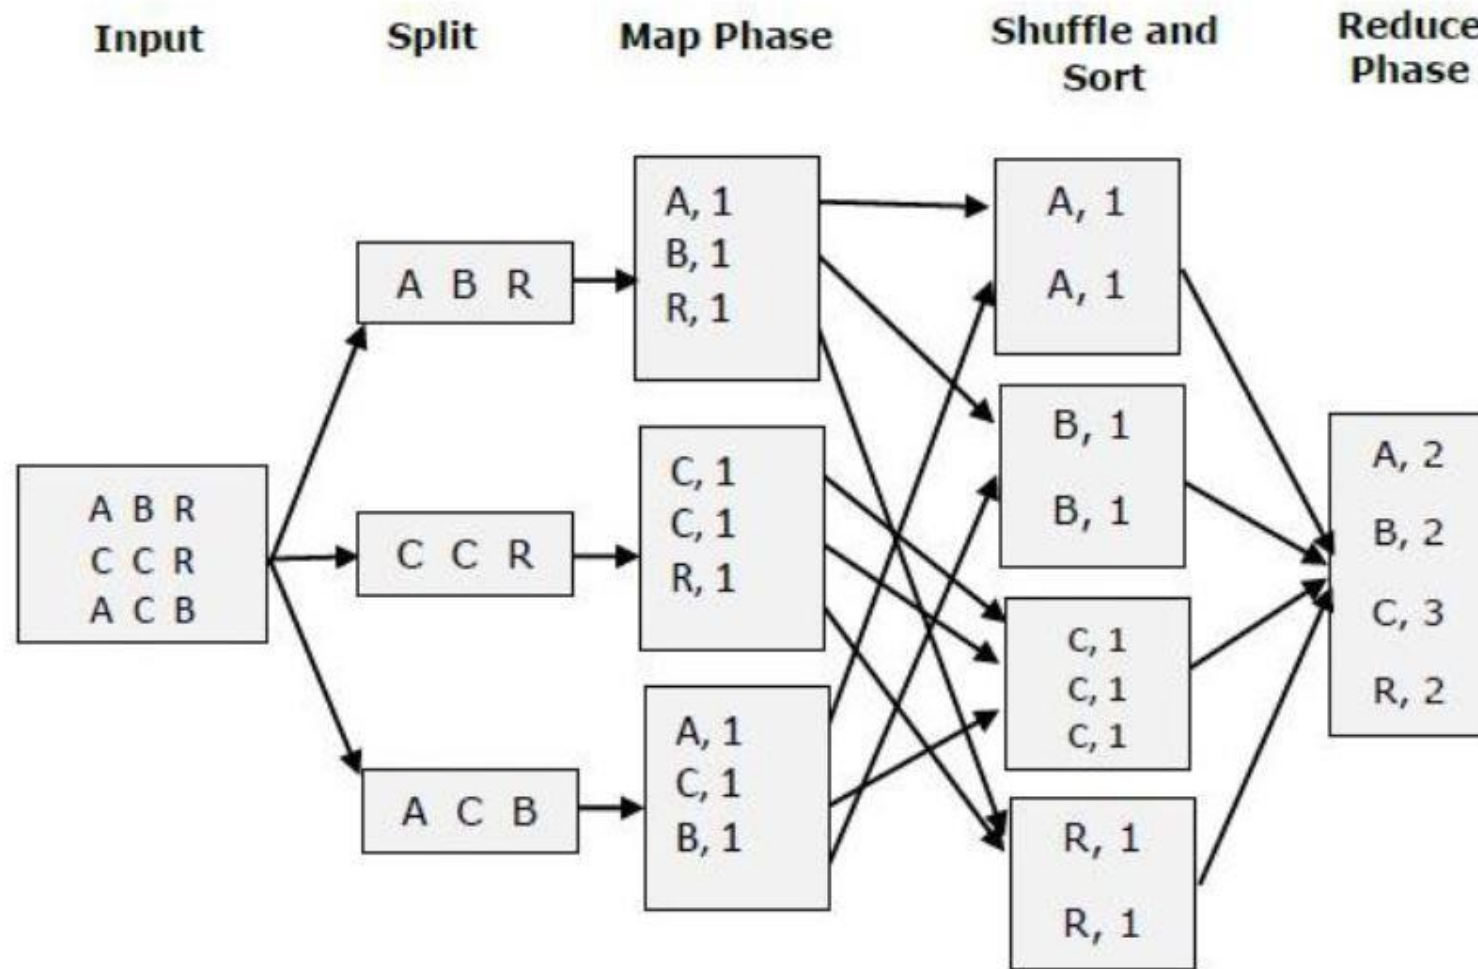


# Phase MapReduce

---

- Input Phase
- Map
- Intermediate Keys
- Combiner
- Shuffle and Sort
- Reducer
- Output Phase

# Ilustrasi proses MapReduce



# Contoh sederhana lain : pseudo code

---

## ■ Counting words in a large set of documents

**map**(string value)

//key: document name

//value: document contents

for each word w in value

*EmitIntermediate*(w, "1");

**reduce**(string key, iterator values)

//key: word

//values: list of counts

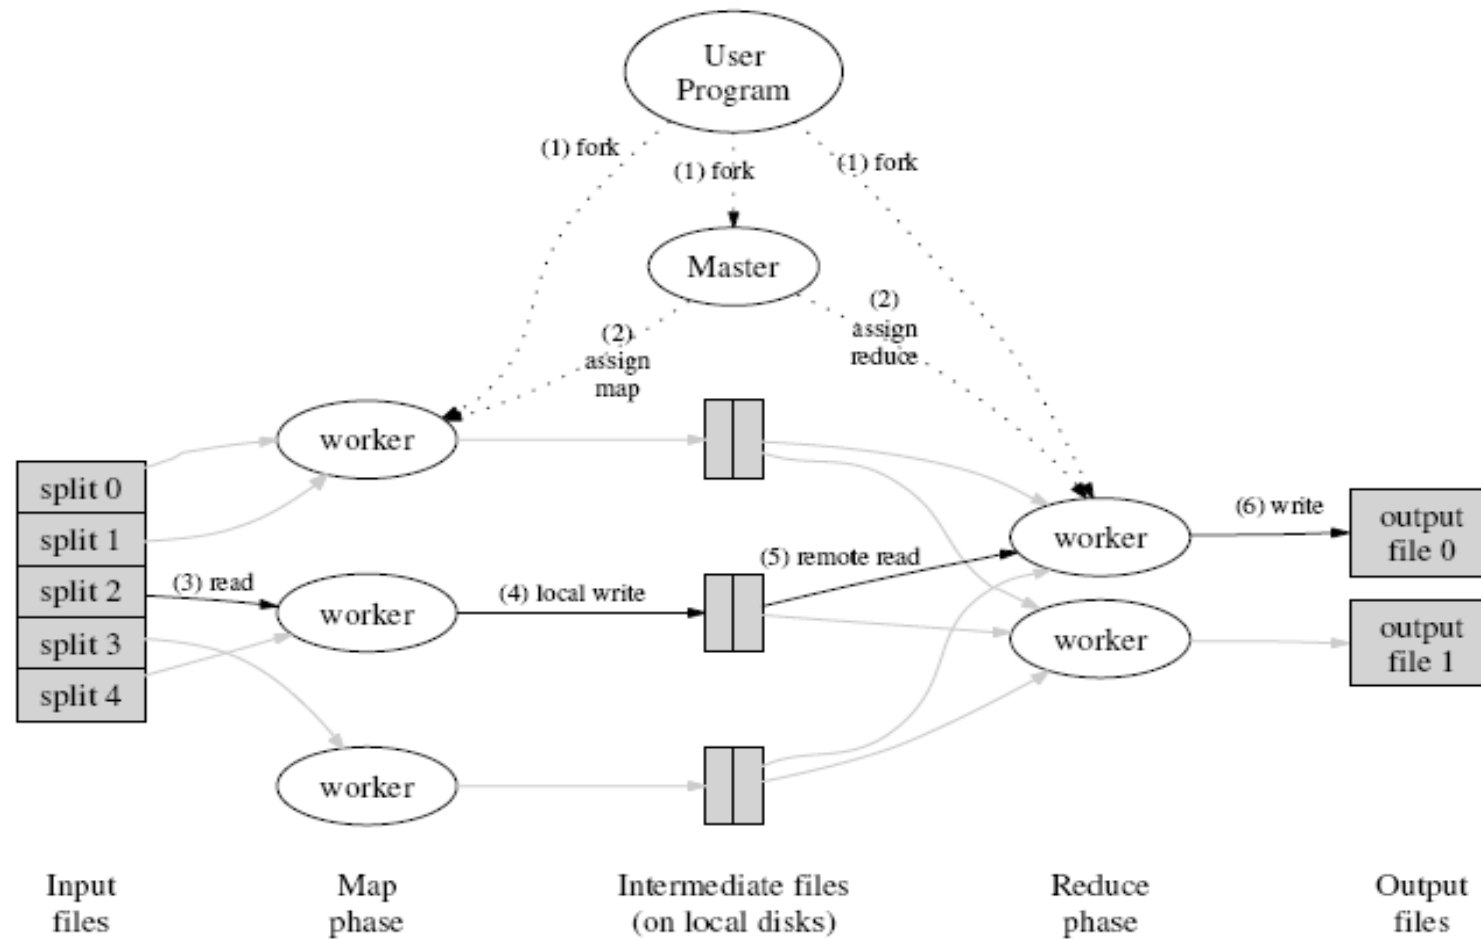
int results = 0;

for each v in values

result += ParseInt(v);

*Emit*(AsString(result));

# Proses Word Count



# Locality issue

---

- Master scheduling policy
  - Asks GFS for locations of replicas of input file blocks
  - Map tasks typically split into 64MB (== GFS block size)
  - Map tasks scheduled so GFS input block replica are on same machine or same rack
- Effect
  - Thousands of machines read input at local disk speed
  - Without this, rack switches limit read rate

# Fault Tolerance

---

- Reactive way
  - Worker failure
    - Heartbeat, Workers are periodically pinged by master
      - NO response = failed worker
    - If the processor of a worker fails, the tasks of that worker are reassigned to another worker.
  - Master failure
    - Master writes periodic checkpoints
    - Another master can be started from the last checkpointed state
    - If eventually the master dies, the job
- Input error: bad records
  - Map/Reduce functions sometimes fail for particular inputs
  - Best solution is to debug & fix, but not always possible
  - On segment fault
    - Send UDP packet to master from signal handler
    - Include sequence number of record being processed
  - Skip bad records
    - If master sees two failures for same record, next worker is told to skip the record


# Lab research Map Reduce

## MapReduce status: MR\_Indexer-beta6-large-2003\_10\_28\_00\_03

Started: Fri Nov 7 09:51:07 2003 -- up 0 hr 37 min 01 sec

1707 workers; 1 deaths

Type	Shards	Done	Active	Input(MB)	Done(MB)	Output(MB)
<a href="#">Map</a>	13853	13853	0	878934.6	878934.6	523499.2
Shuffle	500	500	0	523499.2	520468.6	520468.6
<a href="#">Reduce</a>	500	406	94	520468.6	512265.2	514373.3



### Counters

Variable	Minute	
Mapped (MB/s)	0.0	
Shuffle (MB/s)	0.0	
Output (MB/s)	849.5	
doc-index-hits	0	10
docs-indexed	0	
dups-in-index-merge	0	
mr-merge-calls	35083350	
mr-merge-outputs	35083350	



# Lab Research : Perbaikan2

---

## Refinement


- Task Granularity
  - Minimizes time for fault recovery
  - load balancing
- Local execution for debugging/testing
- Compression of intermediate data

## Notes to emphasized

- No *reduce* can begin until *map* is complete
- Master must communicate locations of intermediate files
- Tasks scheduled based on location of data
- If *map* worker fails any time before *reduce* finishes, task must be completely rerun
- MapReduce library does most of the hard work for us!

# Model Map Reduce : implementasi kasus

## ■ MapReduce Programs In Google Source Tree



Contoh implementasi dalam beberapa kasus lain:

distributed grep

term-vector / host

document clustering

...

distributed sort

web access log stats

machine learning

...

web link-graph reversal

inverted index construction

statistical machine translation

...

# Contoh program word count ( 1 )

---

- Map :

```
#include "mapreduce/mapreduce.h"

// User's map function
class WordCounter : public Mapper {
public:
    virtual void Map(const MapInput& input) {
        const string& text = input.value();
        const int n = text.size();
        for (int i = 0; i < n; ) {
            // Skip past leading whitespace
            while ((i < n) && isspace(text[i]))
                i++;

            // Find word end
            int start = i;
            while ((i < n) && !isspace(text[i]))
                i++;

            if (start < i)
                Emit(text.substr(start, i-start), "1");
        }
    }
};

REGISTER_MAPPER(WordCounter);
```

# Contoh program word count ( 2 )

---

- Reduce :

```
// User's reduce function
class Adder : public Reducer {
    virtual void Reduce(ReduceInput* input) {
        // Iterate over all entries with the
        // same key and add the values
        int64 value = 0;
        while (!input->done()) {
            value += StringToInt(input->value());
            input->NextValue();
        }

        // Emit sum for input->key()
        Emit(IntToString(value));
    }
};
REGISTER_REDUCER(Adder);
```

# Contoh program word count ( 3 )

- Main :

```
int main(int argc, char** argv) {
    ParseCommandLineFlags(argc, argv);

    MapReduceSpecification spec;

    // Store list of input files into "spec"
    for (int i = 1; i < argc; i++) {
        MapReduceInput* input = spec.add_input();
        input->set_format("text");
        input->set_filepattern(argv[i]);
        input->set_mapper_class("WordCounter");
    }

    // Specify the output files:
    //    /gfs/test/freq-00000-of-00100
    //    /gfs/test/freq-00001-of-00100
    //    ...
    MapReduceOutput* out = spec.output();
    out->set_filebase("/gfs/test/freq");
    out->set_num_tasks(100);
    out->set_format("text");
    out->set_reducer_class("Adder");

    // Optional: do partial sums within map
    // tasks to save network bandwidth
    out->set_combiner_class("Adder");

    // Tuning parameters: use at most 2000
    // machines and 100 MB of memory per task
    spec.set_machines(2000);
    spec.set_map_megabytes(100);
    spec.set_reduce_megabytes(100);

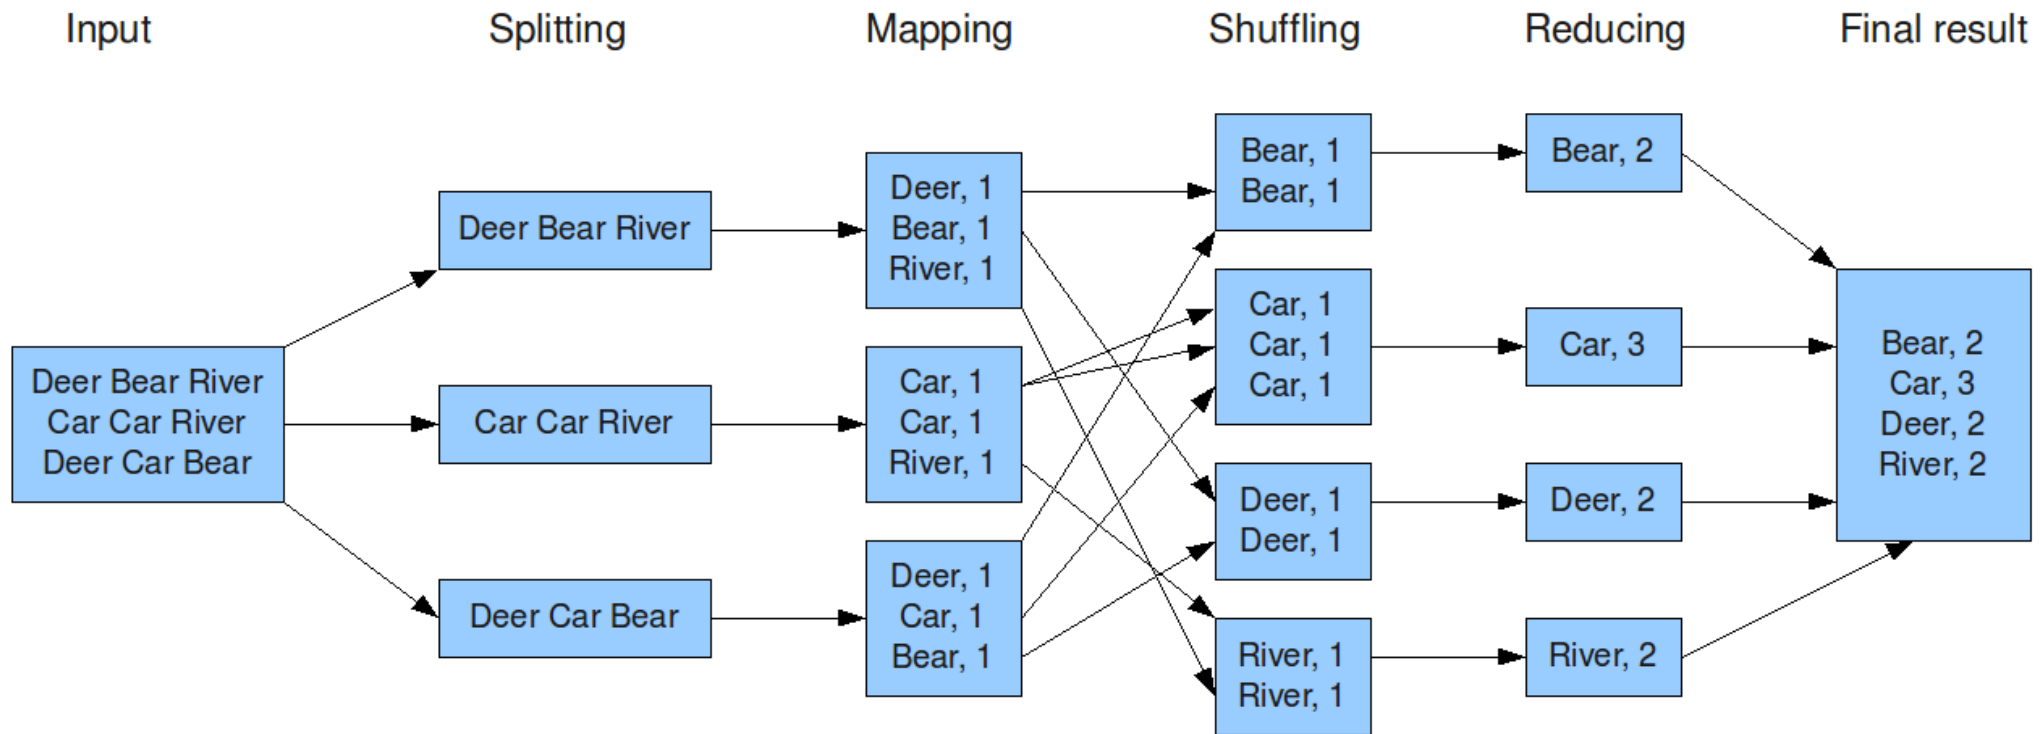
    // Now run it
    MapReduceResult result;
    if (!MapReduce(spec, &result)) abort();

    // Done: 'result' structure contains info
    // about counters, time taken, number of
    // machines used, etc.

    return 0;
}
```

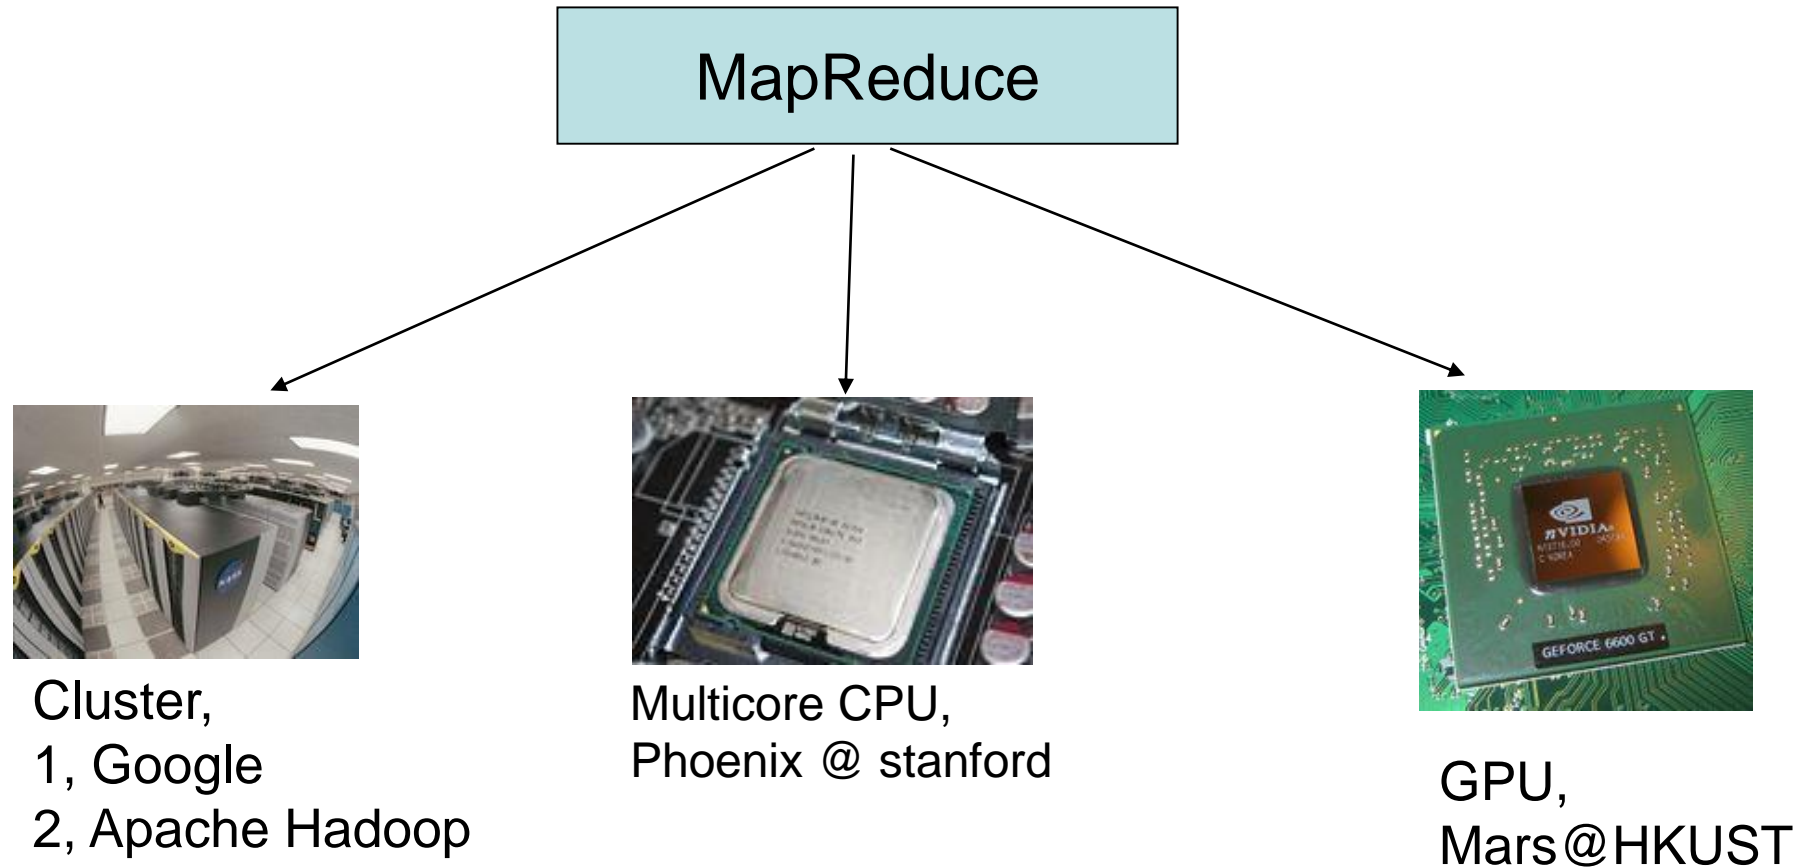
# Map Reduce Process : Word Count

The overall MapReduce word count process



# Implementasi Map Reduce

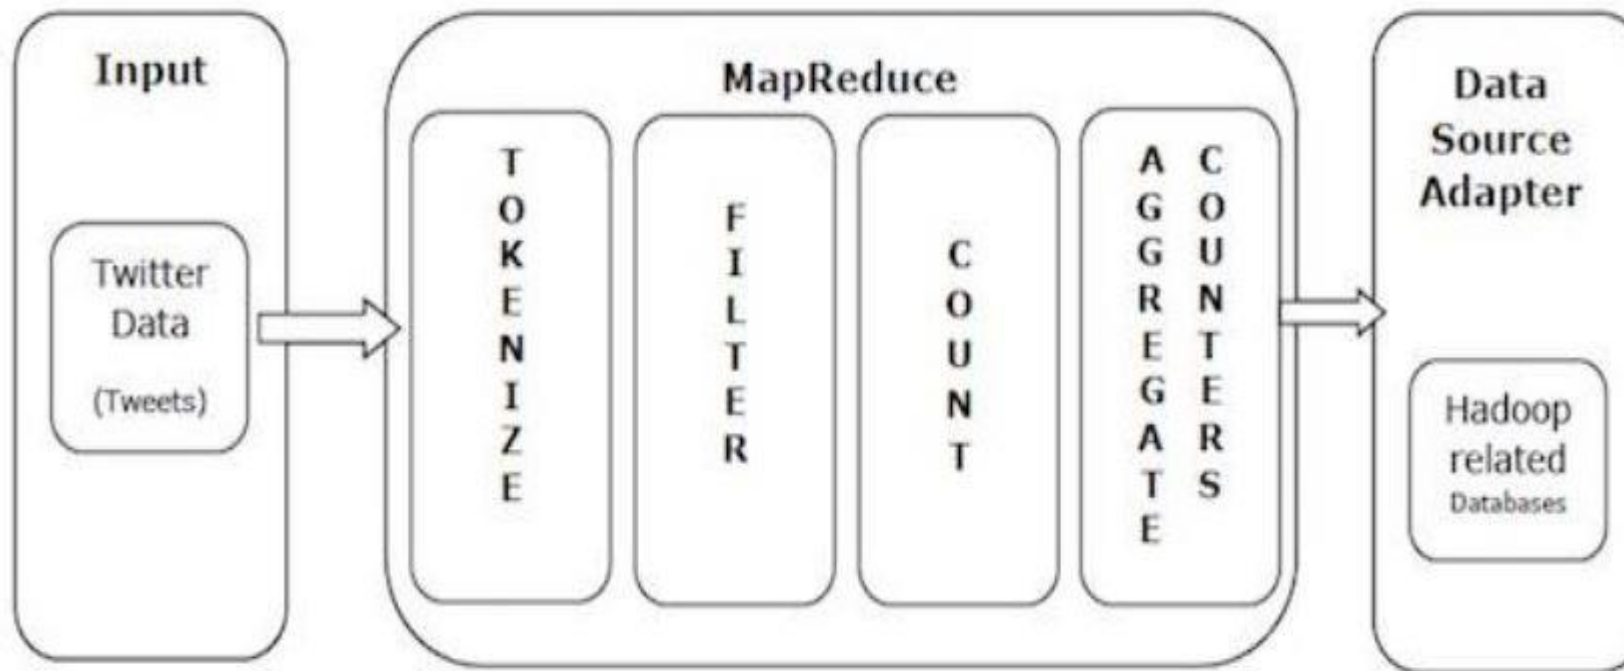
---



# Contoh : MapReduce Twitter

---

- Twitter menerima data setiap hari sekitar 500 juta tweet , 3000 tweet per detik





# Next Hadoop !

---

<b>Google</b>	<b>Yahoo / Apache</b>
<b>MapReduce</b>	<b>Hadoop</b>
<b>GFS</b>	<b>HDFS</b>
<b>Bigtable</b>	<b>HBase</b>
<b>Chubby (distributed lock service)</b>	<b>Apache Zookeeper</b>

# Kuis !!

---

- Jelaskan apa Map Reduce ?
- Mengapa Map Reduce dibutuhkan untuk proses data skala besar ?